

Advanced Docker Best Practices

Kostas Kalevras, DevOps Engineer, GUNet


Agenda

- Part 1
 - Business as Usual
 - Problems
 - The case for Docker
 - Part 2
 - Docker best practices
-




Business as Usual

- 1 VM per service and institutions
 - Service Stack
 - Debian OS (JeOS)
 - Debian packages
 - Custom binaries
 - Configuration files
 - Static/dynamic web pages
 - Installation - Management
 - Manually
 - Scripts
-



Problems

- Imperative
 - Define how to do the task, not where you want to go
 - Hard to automate
 - Hard to scale
 - Development and Operations are very hard to split
 - Service Management
 - VM package management
 - Service binaries
 - Service configuration
-



The case for Docker 1/3

- Lightweight self-contained, isolated mini “VMs”
- Well-defined images (Dockerfile commands). Basically, a Service as Configuration
- Image stacks (every image extends an existing image)
- Image registries
- Version tagging (mariadb:latest, mariadb:10.6)
- Known image size
- UnionFS (shared layers)

The case for Docker 2/3

- Layered Stack
 - Immutable image (service binaries and base configuration)
 - Env variables for per-container configuration
 - Volumes to hold actual data (host filesystem)
- Low-cost container creation/destruction
- Container – VM decoupling – Thin Docker Host
- Defined resource limits (CPU, memory, logs)
- Plugin Interface

The case for Docker 3/3


- Custom images and Docker compose stacks
- GitOps
- Remote Management (Docker contexts)
- Declarative
 - Define the end state
 - Easy to automate
 - Relatively easy to scale

Part 2

Docker best practices

Distributions – Image slim versions

- Use the distribution you know
 - ie Debian, not Alpine
 - You gain in size but use what you already know
- Try using *slim* versions and add any packages you need
 - Bullseye-slim vs bullseye (80MB < 125MB)
 - Python slim vs vanilla python (130MB < 920MB)
 - Node slim vs vanilla node (190MB < 910MB)



```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

Layers


- Each RUN creates a **new** layer in UnionFS
 - You **cannot** delete previous layers
 - The rm below will have no effect

```
RUN wget http://<site>/package.zip  
RUN unzip package.zip  
RUN rm -f package.zip
```

- Run all commands in **one** layer:

```
RUN wget http://<site>/package.zip && unzip package.zip && rm  
-f package.zip
```

- Start with the most basic and less changed layers first (Build caching, shared layers)
 - Install packages
 - Code
 - Configuration files
 - When installing packages always clean-up in the same RUN layer
-



```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

Multi-stage builds

- Don't use large images or install packages only for building
 - Use multi-stage building: A **separate** container is used as builder
 - Example: Use eclipse JDK to build JRE while the base image is *debian:bullseye-slim*
-

```
FROM eclipse-temurin:17-jdk as jre-build
```

```
# Create a custom Java runtime
```

```
RUN $JAVA_HOME/bin/jlink \  
    --add-modules ALL-MODULE-PATH \  
    --strip-debug \  
    --no-man-pages \  
    --no-header-files \  
    --compress=2 \  
    --output /javaruntime
```

```
FROM debian:bullseye-slim
```

```
# Copy Java from OpenJDK
```

```
RUN mkdir -p /javaruntime  
ENV JAVA_HOME=/opt/java/openjdk  
ENV PATH "${JAVA_HOME}/bin:${PATH}"  
COPY --from=jre-build /javaruntime  
$JAVA_HOME
```

```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

Docker entry point

- Create a Docker entrypoint to run various tasks on startup and finally call your application using *exec*
- In the *Dockerfile* you set the **ENTRYPOINT** to point to your entrypoint script and pass arguments using the **CMD** instruction

```
ENTRYPOINT [ "/usr/local/bin/docker-entrypoint.sh" ]  
CMD [ "--ini", "/etc/privayidea/uwsgi/uwsgi.ini" ]
```

- Benefits
 - Signals go to your application, not the shell
 - Setup things on first run
 - Substitute environment vars in configuration files using *envsubst*
 - Decrypt keys using a passphrase passed as env var
-

```
#!/bin/bash  
  
# Create database if it does not exist (in the case of sqlite)  
grep SQLALCHEMY_DATABASE_URI ${PI_CONFIG} | grep -q sqlite && \  
[ ! -f /etc/privacyidea/db/data.sqlite ] && \  
echo "Creating sqlite database.." && ${PI_MANAGE} create_tables  
  
# Decrypt keys  
ENCKEY=$(grep PI_ENCKEY ${PI_CONFIG} | awk '{print $3}' | tr -d \\  
ENCKEY_ENC=$(grep PI_ENCKEY_ENC ${PI_CONFIG} | awk '{print $3}' | tr -d \\  
's/$/.aes/')  
  
PRIVKEY=$(grep PI_PRIVKEY_PRIVATE ${PI_CONFIG} | awk '{print $3}' | tr -d \\  
PRIVKEY_ENC=$(grep PI_PRIVKEY_PRIVATE ${PI_CONFIG} | awk '{print $3}' | tr -d \  
's/$/.pem/')  
  
if [[ -f ${ENCKEY_ENC} ]]; then  
    echo "Decrypting enckey '${ENCKEY_ENC}' to ${ENCKEY} .."  
    openssl enc -aes-256-cbc -md sha512 -pbkdf2 -iter 100000 -d -salt -pass  
pass:${ENCKEY_PASSPHRASE} -in ${ENCKEY_ENC} -out ${ENCKEY}  
fi  
  
if [[ -f ${PRIVKEY_ENC} ]]; then  
    echo "Decrypting private key ${PRIVKEY_ENC} to ${PRIVKEY} .."  
    openssl rsa -passin pass:${ENCKEY_PASSPHRASE} -in ${PRIVKEY_ENC} -out  
${PRIVKEY}  
fi  
  
# Use template to substitute env variables  
envsubst </etc/privacyidea/templates/main.json >/etc/privacyidea/main.json  
  
exec /usr/local/bin/uwsgi "$@"
```



```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

Dependency stack

- In Docker Compose we can setup a dependency stack
 - A service only starts if the dependencies have started
 - We can add a **healthy** requirement

```
depends_on:  
  db:  
    condition: service_healthy  
  proxy:  
    condition: service_healthy
```

- Cron Jobs
 - Docker does not include Cron jobs
 - Kubernetes does
 - We can approximate with a script that sleeps between runs

```
-- docker-compose.rotateaudit.yaml --  
version: '3.7'  
  
services:  
  rotateaudit:  
    image: gunet/privacyidea:${INSTITUTION:-latest}  
    entrypoint: /usr/local/bin/rotate_auditlog.sh  
    environment:  
      - ROTATE_DAYS=120  
      - ROTATE_SLEEP=86400  
    restart: unless-stopped  
    network_mode: "none"  
  
-- rotate_auditlog.sh --  
#!/bin/bash  
  
if [[ ! -v ROTATE_DAYS ]]; then  
  echo "No ROTATE_DAYS env variable available!"  
  exit 1  
fi  
if [[ ! -v ROTATE_SLEEP ]]; then  
  echo "No ROTATE_SLEEP env variable available!"  
  exit 1  
fi  
  
while true  
do  
  echo "Rotating audit log, trimming anything older than ${ROTATE_DAYS}"  
  ${PI_MANAGE} audit rotate_audit --age ${ROTATE_DAYS}  
  echo "Sleeping for ${ROTATE_SLEEP} secs.."  
  sleep ${ROTATE_SLEEP}  
done
```

```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

Dynamic Routing


- Docker Compose stack with multiple services
 - Traefik dynamic routing
 - Configuration using Docker Compose labels
 - Routing
 - HTTP (Host header)
 - HTTPS (SNI)
 - Client-IP
 - HTTP headers (X-Forwarded-For)
 - PROXY2 protocol (SSL)
 - Apache supports both
 - Middleware
 - BasicAuth
 - IP Whitelisting
-

```
guest:  
  image: ghcr.io/gunet/wifi-guest  
  labels:  
    # Explicitly instruct Traefik to expose this service  
    - traefik.enable=true  
    # Router configuration  
    ## One should replace the `guest` name with the name of their service  
    ## Enable TLS and passthrough  
    - traefik.tcp.routers.guest.tls=true  
    - traefik.tcp.routers.guest.tls.passthrough=true  
    ## Listen to the `web` entrypoint  
    - traefik.tcp.routers.guest.entrypoints=https  
    ## Rule based on the HostSNI of the request  
    - traefik.tcp.routers.guest.rule=HostSNI(`www.gunet.gr`)  
    # Service configuration  
    ## 443 is the port that the guest container is listening to  
    - traefik.tcp.services.guest.loadbalancer.server.port=443
```

```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

GitHub integration


- GitHub Container Registry (GHCR)
- Integrate repo with Docker images
- Inherit access rights
 - Private repo has private packages
 - Only users with access to repo can access packages
 - Easy integration in GitHub Actions



```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip  
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

GitHub Actions

- Build images using GitHub actions
 - Build on
 - Push
 - Pull Request
 - Manual invocation
 - Common restrictions
 - Branch
 - Folders/files
 - Integrated with GHCR
 - Workflow stacks
 - Can be used for
 - Tests
 - Linting
-




```
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
RUN wget http://<site>/package.zip RUN unzip package.zip RUN rm -f package.zip
```

GitHub Actions Example

```
name: build-grafana
run-name: Build Grafana Docker image

on:
  workflow_dispatch:
  push:
    branches:
      - "main"
    paths:
      - 'grafana/**'

jobs:
  build-grafana:
    uses: ../.github/workflows/all_imagesbuild_worker.yml
    with:
      image-name: grafana
      directory-name: grafana
      dockerfile: Dockerfile
      secrets: inherit
```

```
jobs:
  build-base:
    runs-on: ubuntu-latest
    strategy:
      matrix:
        institution: [aegean, asfa, aueb, ihu, ionio, latest, staging, unipi, uoi, uop, uowm]
    steps:
      - name: Print inputs
        run: |
          echo "image-name: ${ inputs.image-name }"
          echo "directory-name: ${ inputs.directory-name }"
          echo "institution: ${ matrix.institution }"
          echo "dockerfile ${ inputs.dockerfile }"
      - name: Set timezone
        uses: zcong1993/setup-timezone@master
        with:
          timezone: "Europe/Athens"
      - name: Get date
        id: date
        run: echo "push_date=$(date '+%H:%M@%d/%m/%Y')" >> $GITHUB_ENV
      - name: Checkout
        uses: actions/checkout@v3
      - name: Login to GitHub Container Registry
        uses: docker/login-action@v2
        with:
          registry: ghcr.io
          username: ${ github.actor }
          password: ${ secrets.GITHUB_TOKEN }
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v2
      - name: Build and push Base image
        uses: docker/build-push-action@v3
        with:
          context: ../${ inputs.directory-name }
          file: ../${ inputs.directory-name }/${ inputs.dockerfile }
          build-args: |
            INSTITUTION=${ matrix.institution }
          push: true
          tags: ghcr.io/gunet/${ inputs.image-name }-${ matrix.institution }
          labels: gr.gunet.${ inputs.image-name }.pushdate=${ env.push_date }
          cache-from: type=gha
          cache-to: type=gha,mode=max
```

Thank you

Comments - Questions